

```
from google.colab import files
import pandas as pd

files.upload()
df=pd.read_csv('climate_change_dataset.csv')
print(df)
```

 Choose Files

No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving climate\_change\_dataset.csv to climate\_change\_dataset.csv

	Year	Country	Avg Temperature (°C)	C02 Emissions (Tons/Capita)	\
0	2006	UK	8.9	9.3	
1	2019	USA	31.0	4.8	
2	2014	France	33.9	2.8	
3	2010	Argentina	5.9	1.8	
4	2007	Germany	26.9	5.6	
..	...	...	...	...	
995	2019	India	23.6	8.0	
996	2000	UK	21.8	10.0	
997	2019	Argentina	23.8	8.9	
998	2016	Australia	21.0	14.9	
999	2011	Germany	24.1	17.3	

	Sea Level Rise (mm)	Rainfall (mm)	Population	Renewable Energy (%)	\
0	3.1	1441	530911230	20.4	
1	4.2	2407	107364344	49.2	
2	2.2	1241	441101758	33.3	
3	3.2	1892	1069669579	23.7	
4	2.4	1743	124079175	12.5	
..	...	...	...	...	
995	1.2	1365	1358019778	10.0	
996	2.2	1273	876123161	14.9	
997	4.7	891	1120533308	25.9	
998	3.1	1136	380662109	24.5	
999	2.1	2854	398407112	41.0	

	Extreme Weather Events	Forest Area (%)
0	14	59.8
1	8	31.0
2	9	35.5
3	7	17.7
4	4	17.4
..	...	...
995	8	20.2
996	14	30.1
997	10	46.5
998	3	44.5
999	3	19.8

[1000 rows x 10 columns]

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

Basic Info

```
print("Basic Information:")
print(df.info())
```

```
Basic Information:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Year                                1000 non-null   int64
1   Country                            1000 non-null   object
2   Avg Temperature (°C)               1000 non-null   float64
3   CO2 Emissions (Tons/Capita)        1000 non-null   float64
4   Sea Level Rise (mm)                1000 non-null   float64
5   Rainfall (mm)                      1000 non-null   int64
6   Population                          1000 non-null   int64
7   Renewable Energy (%)               1000 non-null   float64
8   Extreme Weather Events             1000 non-null   int64
9   Forest Area (%)                    1000 non-null   float64
dtypes: float64(5), int64(4), object(1)
memory usage: 78.3+ KB
None
```

INFERENCE

The dataset contains 1000 records with 10 columns related to climate and environmental factors.

Summary Statistics

```
print("\nSummary Statistics:")
print(df.describe(include='all'))
```

```
Summary Statistics:
      Year  Country  Avg Temperature (°C) \
count  1000.000000    1000    1000.000000
unique      NaN        15             NaN
top         NaN  Indonesia             NaN
freq         NaN         75             NaN
mean    2011.432000      NaN        19.883100
std         7.147199      NaN        8.542897
min     2000.000000      NaN        5.000000
25%     2005.000000      NaN        12.175000
50%     2012.000000      NaN        20.100000
75%     2018.000000      NaN        27.225000
max     2023.000000      NaN        34.900000

      CO2 Emissions (Tons/Capita)  Sea Level Rise (mm)  Rainfall (mm) \
count      1000.000000      1000.000000      1000.000000
unique           NaN           NaN           NaN
```

top	NaN	NaN	NaN
freq	NaN	NaN	NaN
mean	10.425800	3.009600	1738.761000
std	5.614665	1.146081	708.976616
min	0.500000	1.000000	501.000000
25%	5.575000	2.000000	1098.750000
50%	10.700000	3.000000	1726.000000
75%	15.400000	4.000000	2362.500000
max	20.000000	5.000000	2999.000000

	Population	Renewable Energy (%)	Extreme Weather Events \
count	1.000000e+03	1000.000000	1000.000000
unique	NaN	NaN	NaN
top	NaN	NaN	NaN
freq	NaN	NaN	NaN
mean	7.053830e+08	27.300500	7.291000
std	4.093910e+08	12.970808	4.422655
min	3.660891e+06	5.100000	0.000000
25%	3.436242e+08	16.100000	3.000000
50%	7.131166e+08	27.150000	8.000000
75%	1.073868e+09	38.925000	11.000000
max	1.397016e+09	50.000000	14.000000


	Forest Area (%)
count	1000.000000
unique	NaN
top	NaN
freq	NaN
mean	40.572000
std	17.398998
min	10.100000
25%	25.600000
50%	41.150000
75%	55.800000
max	70.000000

# INFERENCE

The dataset provides summary statistics for 1000 records, covering climate-related factors. The average temperature is 19.88°C, CO<sub>2</sub> emissions are 10.42 tons per capita, and sea level rise is 3.01 mm on average. The dataset includes population data, renewable energy use, and extreme weather events, with Indonesia being the most frequent country.

## ⌵ Check for Missing Values

```
print("\nMissing Values:")
print(df.isnull().sum())
```



```

Missing Values:
Year                0
Country            0
Avg Temperature (°C) 0
CO2 Emissions (Tons/Capita) 0

```

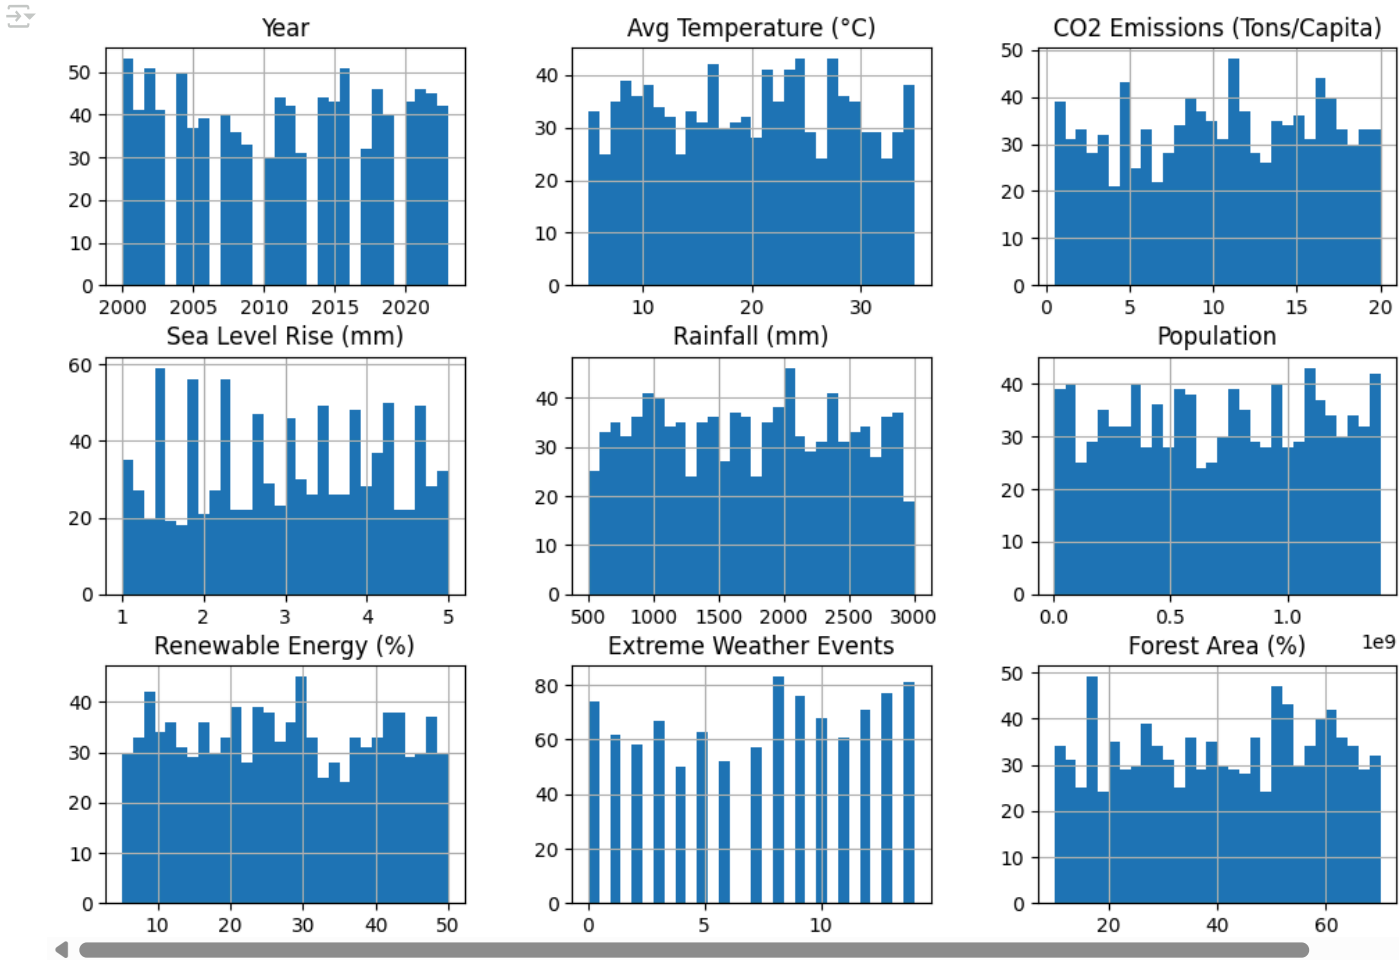
```
Sea Level Rise (mm)      0
Rainfall (mm)            0
Population               0
Renewable Energy (%)     0
Extreme Weather Events   0
Forest Area (%)          0
dtype: int64
```

## INFERENCE

The dataset has no missing values, as all columns contain 0 null entries. This ensures data completeness, making it suitable for analysis without the need for imputation or data cleaning.

### ✦ Distribution of Numerical Features

```
df.hist(figsize=(12, 8), bins=30)
plt.show()
```

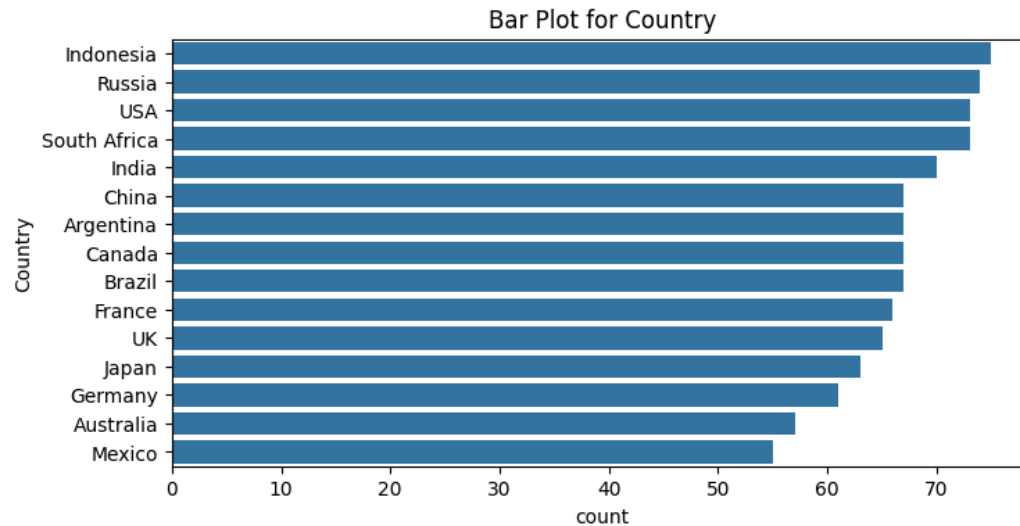


## INFERENCE

The histograms show the distribution of climate-related data. Most variables are spread out, while some, like Sea Level Rise and Extreme Weather Events, have values concentrated in specific ranges. This helps in understanding data trends and patterns.

### ✓ Bar Chart

```
categorical_cols = df.select_dtypes(include=['object']).columns
for col in categorical_cols:
    plt.figure(figsize=(8, 4))
    sns.countplot(y=df[col], order=df[col].value_counts().index)
    plt.title(f'Bar Plot for {col}')
    plt.show()
```




## INFERENCE

The bar plot shows the count of data entries for each country. Indonesia, Russia, and the USA have the highest counts, while other countries have fewer entries. This helps in understanding the distribution of categorical data.

### ✓ Correlation Matrix

```
numerical_df = df.select_dtypes(include=np.number)
plt.figure(figsize=(10, 6))
sns.heatmap(numerical_df.corr(), annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Matrix')
plt.show()
```



## INFERENCE

The correlation matrix visually represents the relationships between different numerical variables. The values range from -1 to 1, where:

1 (red) indicates a perfect positive correlation.

-1 (blue) indicates a perfect negative correlation.

0 (dark blue) means no correlation.

From the heatmap, it seems that most variables have weak correlations with each other. This suggests that they are mostly independent, except for some minor relationships.


## ✓ Boxplots

```
numerical_df = df.select_dtypes(include=np.number)
numerical_cols = numerical_df.columns
for col in numerical_cols:
    plt.figure(figsize=(8, 4))
    sns.boxplot(x=df[col])
    plt.title(f'Boxplot for {col}')
    plt.show()
```







Boxplot for Year




Boxplot for Avg Temperature (°C)

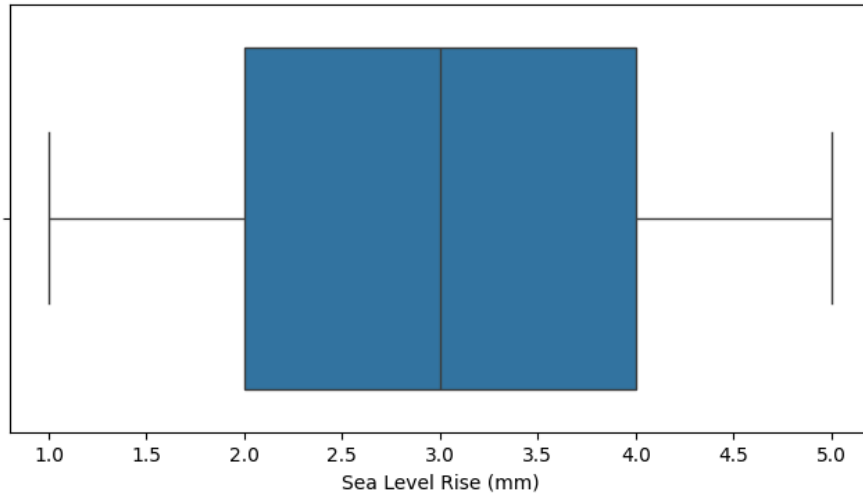


Boxplot for CO2 Emissions (Tons/Capita)

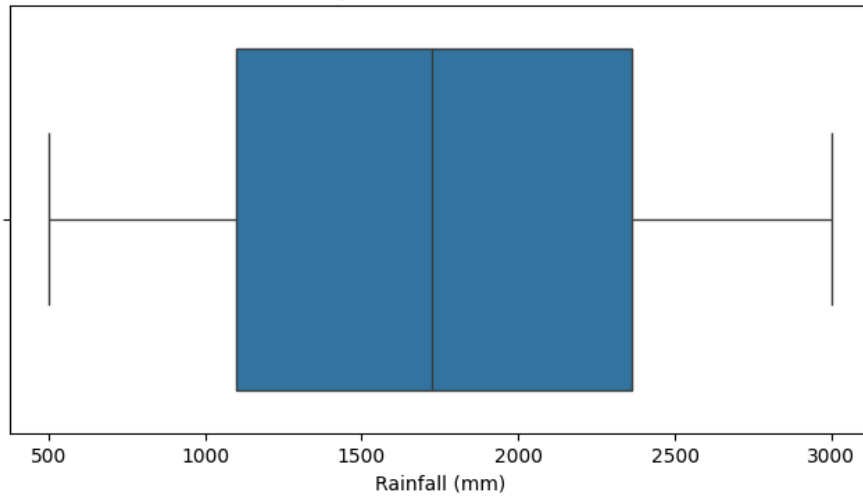





Boxplot for Sea Level Rise (mm)




Boxplot for Rainfall (mm)

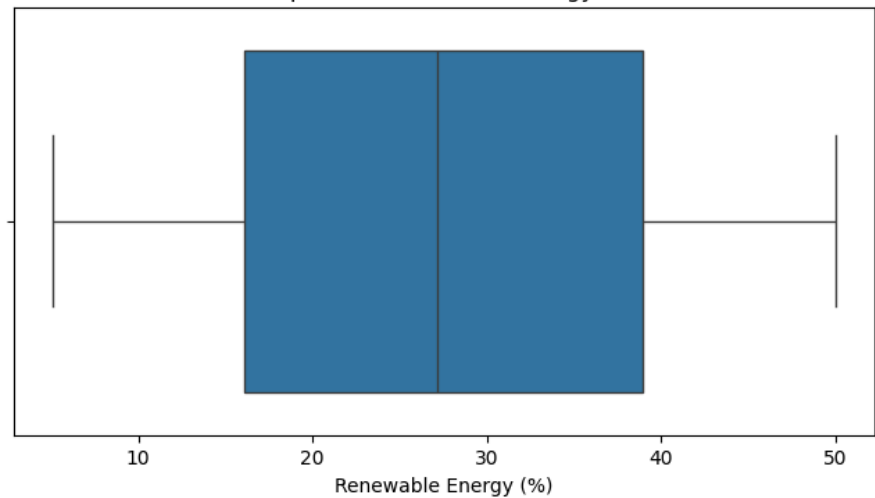


Boxplot for Population

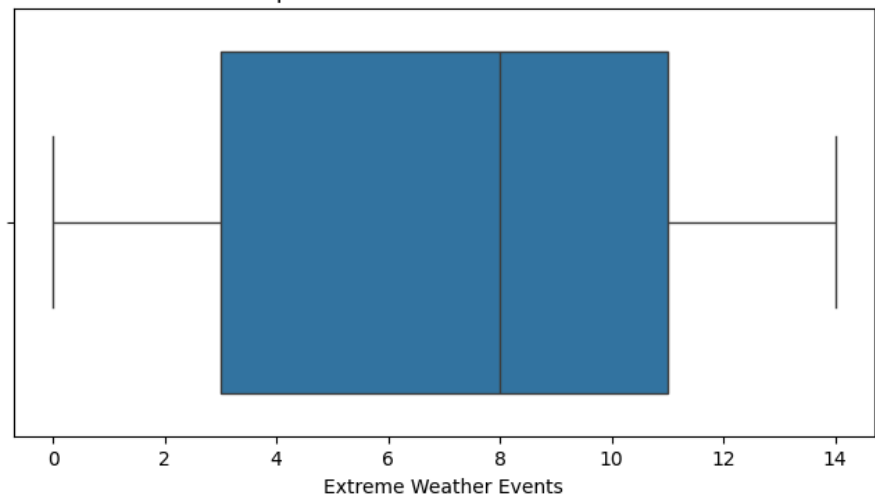




Boxplot for Renewable Energy (%)




Boxplot for Extreme Weather Events



Boxplot for Forest Area (%)





# INFERENCE

The boxplots visualize the distribution of numerical variables, highlighting their median, quartiles, and potential outliers. They help in identifying:

Skewness: If the median is not centered.

Outliers: Points outside the whiskers.


Spread of Data: The interquartile range (IQR).

## ✓ KDE Plots for Distribution Analysis


```
for col in numerical_cols:
    plt.figure(figsize=(8, 4))
    sns.kdeplot(df[col], fill=True)
    plt.title(f'Density Plot for {col}')
    plt.show()
```




Density Plot for Year




Density Plot for Avg Temperature (°C)

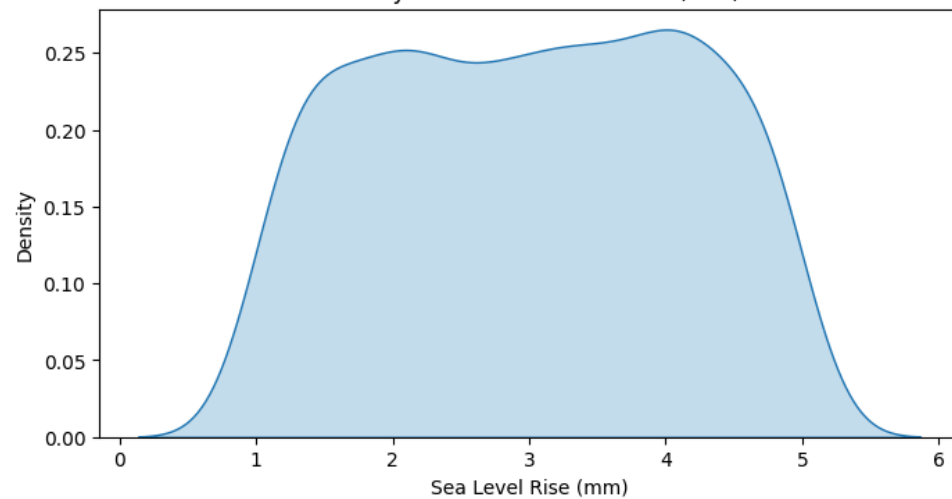


Density Plot for CO2 Emissions (Tons/Capita)







Density Plot for Sea Level Rise (mm)


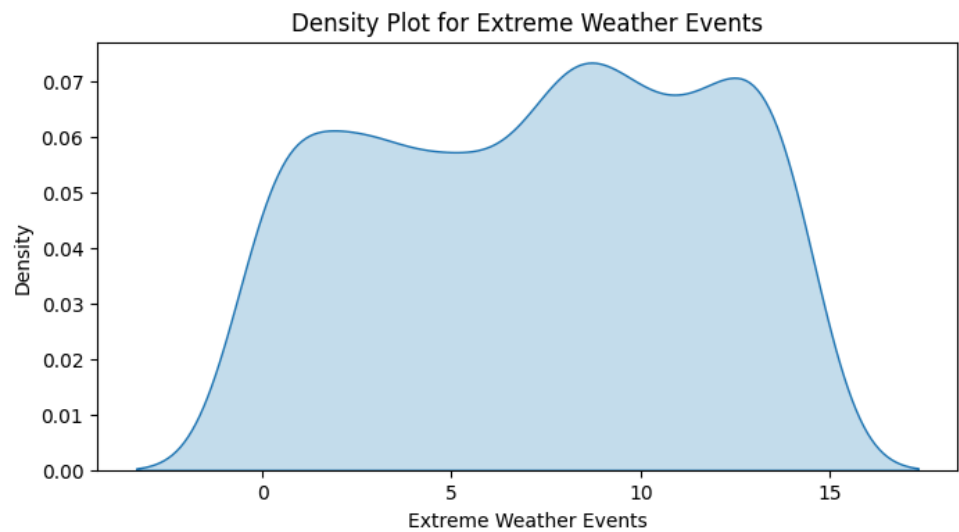
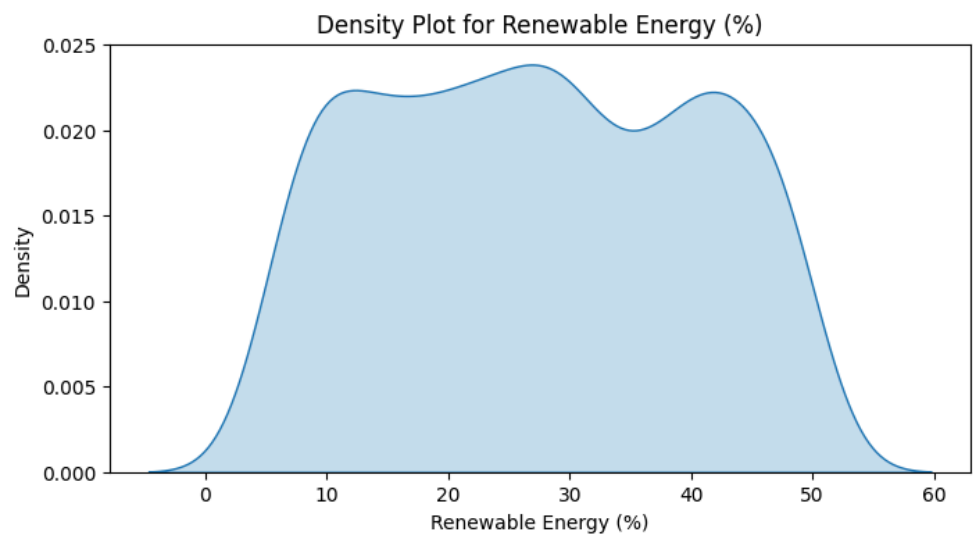



Density Plot for Rainfall (mm)

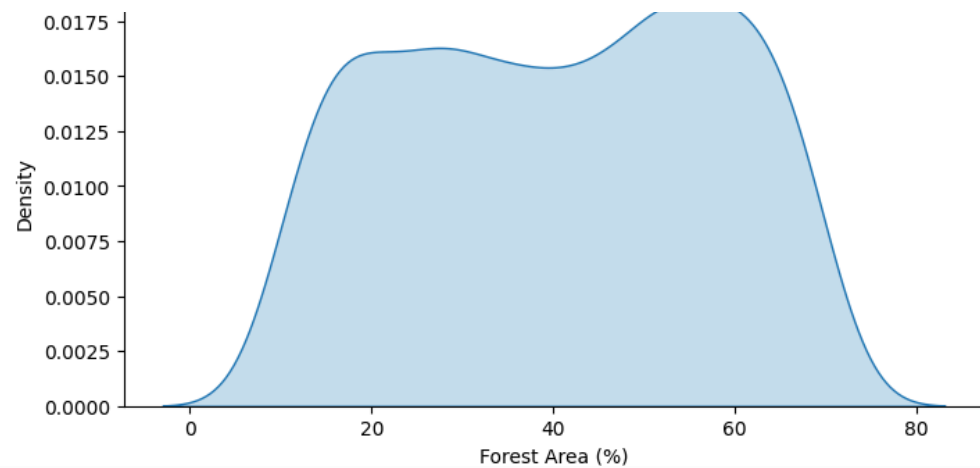


Density Plot for Population









# INFERENCE

The density plots provide insights into the distribution of numerical data. They help in understanding:

Peaks (Modes): Where the data is concentrated.

Skewness: If the distribution is left or right-skewed.

Spread: How dispersed the values are.

## > Line Plots for Trend Analysis

[ ] ↪ 1 cell hidden

# INFERENCE

The line plots show high fluctuations, indicating noisy data with no clear trend. The data varies significantly, suggesting potential randomness or high variability. Further analysis may be needed to identify patterns.

## > Scatter Plots

[ ] ↪ 1 cell hidden

# INFERENCE

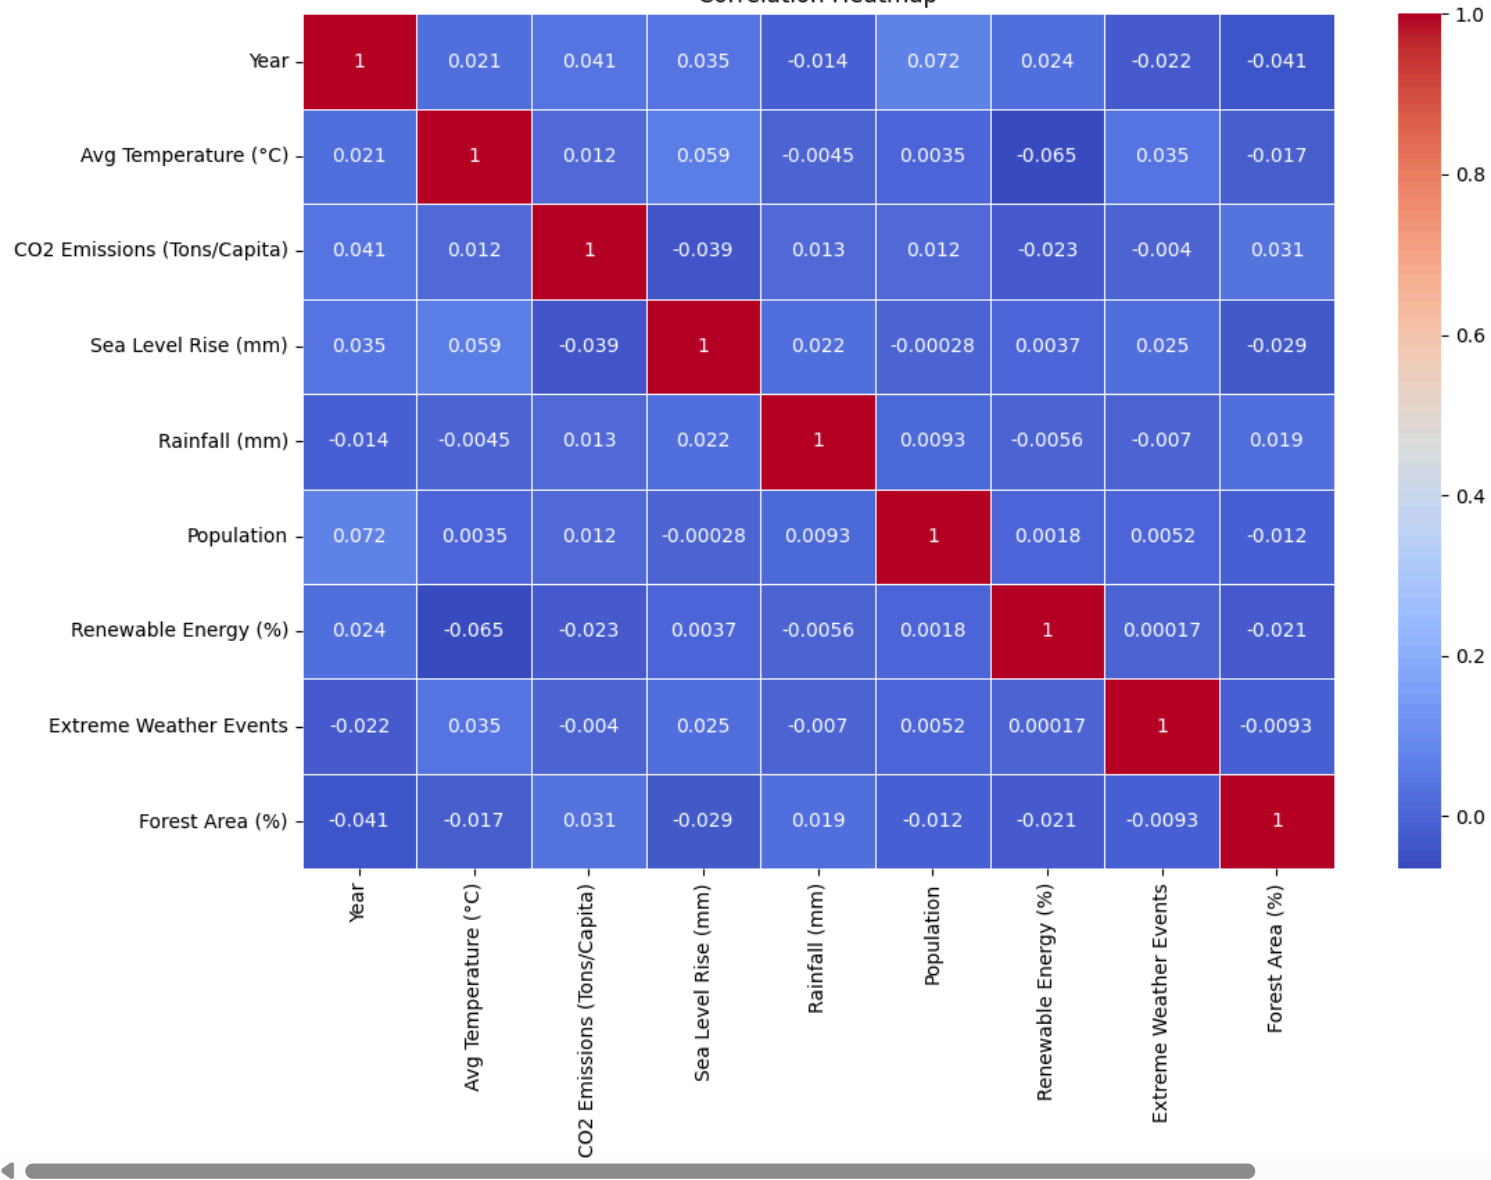
The scatter plots show no clear increasing or decreasing trends over the years. All variables exhibit fluctuations, indicating inconsistent patterns without strong correlations to time.

## ✓ Heatmap for Feature Relationships

```
plt.figure(figsize=(12, 8))
numerical_df = df.select_dtypes(include=np.number)
sns.heatmap(numerical_df.corr(), annot=True, cmap='coolwarm', linewidths=0.5)
plt.title('Correlation Heatmap')
plt.show()
```



Correlation Heatmap



## INFERENCE

The correlation heatmap shows weak correlations among the variables, with no strong relationships observed. The highest correlations are self-correlations (value = 1), while other values remain close to zero, indicating minimal linear dependency between factors.

> Pie Charts for Categorical Data


↳ 1 cell hidden

INFERENCE

The pie chart shows the distribution of countries in the dataset. The proportions are fairly balanced, with some countries having slightly higher representation. No single country dominates, indicating a diverse dataset.

✓ Stem and Leaf Plot

```
from collections import Counter
col = 'Year'
counts = Counter(df[col].astype(int))
print(f"\nStem-and-Leaf Plot for {col}:")
for key in sorted(counts.keys()):
    print(f"{key} | {'*' * counts[key]}")
```



INFERENCE

The stem-and-leaf plot shows the distribution of data points across different years. The years with more occurrences have longer bars of asterisks, indicating higher frequencies. This helps visualize trends or peaks in the dataset over time.

## ✓ Non-Graphical Measures (Central Tendency, Dispersion, Shape)

```
print("\nMean:")
numerical_df = df.select_dtypes(include=np.number)
print(numerical_df.mean())

print("\nMedian:")
numerical_df = df.select_dtypes(include=np.number)
print(numerical_df.median())

print("\nMode:")
print(df.mode().iloc[0])

print("\nVariance:")
numerical_df = df.select_dtypes(include=np.number)
print(numerical_df.var())

print("\nStandard Deviation:")
numerical_df = df.select_dtypes(include=np.number)
print(numerical_df.std())

print("\nRange:")
numerical_df = df.select_dtypes(include=np.number)
print(numerical_df.max() - numerical_df.min())

print("\nSkewness:")
numerical_df = df.select_dtypes(include=np.number)
print(numerical_df.skew())

print("\nKurtosis:")
numerical_df = df.select_dtypes(include=np.number)
print(numerical_df.kurt())
```



```
Mean:
Year                2.011432e+03
Avg Temperature (°C) 1.988310e+01
CO2 Emissions (Tons/Capita) 1.042580e+01
Sea Level Rise (mm)  3.009600e+00
Rainfall (mm)        1.738761e+03
Population            7.053830e+08
Renewable Energy (%)  2.730050e+01
Extreme Weather Events 7.291000e+00
Forest Area (%)       4.057200e+01
dtype: float64
```

```
Median:
Year                2.012000e+03
Avg Temperature (°C) 2.010000e+01
CO2 Emissions (Tons/Capita) 1.070000e+01
Sea Level Rise (mm)  3.000000e+00
Rainfall (mm)        1.726000e+03
Population            7.131166e+08
Renewable Energy (%)  2.715000e+01
```



```
Extreme Weather Events      8.000000e+00
Forest Area (%)             4.115000e+01
dtype: float64

Mode:
Year                2000.0
Country            Indonesia
Avg Temperature (°C)      11.8
CO2 Emissions (Tons/Capita)  11.2
Sea Level Rise (mm)       4.1
Rainfall (mm)           1962.0
Population           3660891
Renewable Energy (%)      20.2
Extreme Weather Events      8.0
Forest Area (%)           27.9
Name: 0, dtype: object

Variance:
Year                5.108246e+01
Avg Temperature (°C)  7.298109e+01
CO2 Emissions (Tons/Capita)  3.152446e+01
Sea Level Rise (mm)    1.313501e+00
Rainfall (mm)         5.026478e+05
Population           1.676010e+17
Renewable Energy (%)   1.682419e+02
Extreme Weather Events  1.955988e+01
Forest Area (%)       3.027251e+02
dtype: float64

Standard Deviation:
Year                7.147199e+00
Avg Temperature (°C)  8.542897e+00
CO2 Emissions (Tons/Capita)  5.614665e+00
Sea Level Rise (mm)    1.146081e+00
Rainfall (mm)         7.089766e+02
Population           4.093910e+08
```

INFERENCE

The statistical summary provides key insights into the dataset, including measures of central tendency (mean, median, mode) and dispersion (variance, standard deviation, range). Skewness and kurtosis help analyze the distribution shape. These metrics highlight trends and variations in environmental and population-related factors.

Skewness: Measures the Asymmetry of Data Distribution

- Skewness  $\approx 0 \rightarrow$  The data is symmetrical (normal distribution).
- Skewness  $> 0 \rightarrow$  The data is right-skewed (longer tail on the right).
- Skewness  $< 0 \rightarrow$  The data is left-skewed (longer tail on the left).

Kurtosis: Measures the "Peakedness" of Distribution

Kurtosis  $\approx 3 \rightarrow$  Normal distribution (mesokurtic).

Kurtosis < 3 → Flat distribution (platykurtic, meaning fewer outliers).

Kurtosis > 3 → Peaked distribution (leptokurtic, meaning many outliers).

## ✓ Value Counts for categorical features

```
for col in categorical_cols:
    print(f"\nValue Counts for {col}:")
    print(df[col].value_counts())
```



```
Value Counts for Country:
Country
Indonesia      75
Russia          74
USA             73
South Africa    73
India           70
China           67
Argentina       67
Canada          67
Brazil          67
France          66
UK              65
Japan           63
Germany         61
Australia       57
Mexico          55
Name: count, dtype: int64
```

## INFERENCE

The value counts show the frequency of records for each country in the dataset. Indonesia has the highest count (75), followed closely by Russia and the USA. This indicates a relatively balanced distribution across countries, with some having slightly more representation.

## ✓ Multivariate Non-Graphical Analysis

```
# Z-Score Calculation
print("\nZ-Scores:")
numerical_df = df.select_dtypes(include=np.number)
z_scores = (numerical_df - numerical_df.mean()) / numerical_df.std()
print(z_scores.head())
```

```
# Covariance Matrix
print("\nCovariance Matrix:")
print(numerical_df.cov())
```

```
# Correlation Matrix
print("\nCorrelation Matrix:")
```

```
numerical_df = df.select_dtypes(include=np.number)
print(numerical_df.corr())
```



Z-Scores:

	Year	Avg Temperature (°C)	CO2 Emissions (Tons/Capita)	\
0	-0.760018	-1.285641	-0.200511	
1	1.058876	1.301303	-1.001983	
2	0.359302	1.640767	-1.358193	
3	-0.200358	-1.636810	-1.536298	
4	-0.620103	0.821372	-0.859499	

	Sea Level Rise (mm)	Rainfall (mm)	Population	Renewable Energy (%)	\
0	0.078878	-0.419987	-0.426174	-0.532002	
1	1.038670	0.942540	-1.460752	1.688368	
2	-0.706407	-0.702084	-0.645547	0.462539	
3	0.166131	0.216141	0.889825	-0.277585	
4	-0.531900	0.005979	-1.419923	-1.141062	

	Extreme Weather Events	Forest Area (%)
0	1.516962	1.105121
1	0.160311	-0.550147
2	0.386419	-0.291511
3	-0.065798	-1.314558
4	-0.744123	-1.331801

Covariance Matrix:

	Year	Avg Temperature (°C)	\
Year	5.108246e+01	1.279380e+00	
Avg Temperature (°C)	1.279380e+00	7.298109e+01	
CO2 Emissions (Tons/Capita)	1.646000e+00	5.910971e-01	
Sea Level Rise (mm)	2.902430e-01	5.777099e-01	
Rainfall (mm)	-6.910886e+01	-2.745019e+01	
Population	2.099217e+08	1.241434e+07	
Renewable Energy (%)	2.208793e+00	-7.254826e+00	
Extreme Weather Events	-6.833954e-01	1.320538e+00	
Forest Area (%)	-5.139143e+00	-2.530924e+00	

	CO2 Emissions (Tons/Capita)	Sea Level Rise (mm)	\
Year	1.646000e+00	0.290243	
Avg Temperature (°C)	5.910971e-01	0.577710	
CO2 Emissions (Tons/Capita)	3.152446e+01	-0.249767	
Sea Level Rise (mm)	-2.497674e-01	1.313501	
Rainfall (mm)	5.296643e+01	17.960555	
Population	2.650091e+07	-129901.767753	
Renewable Energy (%)	-1.700884e+00	0.054340	
Extreme Weather Events	-9.980761e-02	0.126733	
Forest Area (%)	3.066199e+00	-0.571693	

	Rainfall (mm)	Population	\
Year	-6.910886e+01	2.099217e+08	
Avg Temperature (°C)	-2.745019e+01	1.241434e+07	
CO2 Emissions (Tons/Capita)	5.296643e+01	2.650091e+07	
Sea Level Rise (mm)	1.796055e+01	-1.299018e+05	
Rainfall (mm)	5.026478e+05	2.686083e+09	
Population	2.686083e+09	1.676010e+17	
Renewable Energy (%)	-5.187035e+01	9.480119e+06	
Extreme Weather Events	-2.184430e+01	9.369384e+06	
Forest Area (%)	2.290809e+02	-8.290061e+07	

Renewable Energy (%) Extreme Weather Events \



## ✓ INFERENCE

The Z-score standardizes the dataset, highlighting how each value deviates from the mean. The covariance and correlation matrices reveal relationships between variables, showing the strength and direction of their dependencies. This helps in understanding trends and potential predictive factors.

```
import numpy as np

# Z-score calculation
def zscore_outliers(df, threshold=3):
    outlier_indices = []

    for col in numerical_cols:
        mean = np.mean(df[col])
        std = np.std(df[col])

        # Z-score calculation
        z_scores = (df[col] - mean) / std
        outliers = df[np.abs(z_scores) > threshold].index
        outlier_indices.extend(outliers)

    print(f"Feature: {col} | Outliers: {len(outliers)}")

    return list(set(outlier_indices))

# Detect outliers
z_outliers = zscore_outliers(df)
print(f"Total Outliers Detected by Z-score: {len(z_outliers)}")
```

```
→ Feature: Year | Outliers: 0
Feature: Avg Temperature (°C) | Outliers: 0
Feature: CO2 Emissions (Tons/Capita) | Outliers: 0
Feature: Sea Level Rise (mm) | Outliers: 0
Feature: Rainfall (mm) | Outliers: 0
Feature: Population | Outliers: 0
Feature: Renewable Energy (%) | Outliers: 0
Feature: Extreme Weather Events | Outliers: 0
Feature: Forest Area (%) | Outliers: 0
Total Outliers Detected by Z-score: 0
```

No outliers were detected in the dataset based on the Z-score method.

This indicates that the data points fall within a reasonable range of values.

The dataset may already be well-cleaned and standardized, with no extreme anomalies.

Possible reasons for not detecting outliers:

1. Standardized or Preprocessed Data:

The dataset may already be standardized or normalized, bringing all features into a consistent range.

As a result, no data points fall beyond the Z-score threshold.

## 2. Uniform Distribution:

The dataset may have a naturally uniform distribution without extreme values.

This could be the case for datasets with smooth, consistent trends over time (e.g., temperature or population data).

## 3. Z-Score Threshold Setting:

The Z-score threshold was likely set to 3 (a common practice), which means it only identifies extremely distant values as outliers.

If the data is relatively consistent, no values may exceed this threshold.

## 4. Limited Data Range:

If the dataset contains values within a narrow range, it may lack true outliers.

This is common in datasets with values changing gradually over time (e.g., gradual temperature rise).

```
# IQR calculation
def iqr_outliers(df):
    outlier_indices = []

    for col in numerical_cols:
        Q1 = df[col].quantile(0.25)
        Q3 = df[col].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR

        # Detecting outliers
        outliers = df[(df[col] < lower_bound) | (df[col] > upper_bound)].index
        outlier_indices.extend(outliers)

    print(f"Feature: {col} | Outliers: {len(outliers)}")

    return list(set(outlier_indices))

# Detect outliers
iqr_outliers = iqr_outliers(df)
print(f"Total Outliers Detected by IQR: {len(iqr_outliers)}")
```

```
🔗 Feature: Year | Outliers: 0
Feature: Avg Temperature (°C) | Outliers: 0
Feature: CO2 Emissions (Tons/Capita) | Outliers: 0
Feature: Sea Level Rise (mm) | Outliers: 0
Feature: Rainfall (mm) | Outliers: 0
Feature: Population | Outliers: 0
Feature: Renewable Energy (%) | Outliers: 0
Feature: Extreme Weather Events | Outliers: 0
Feature: Forest Area (%) | Outliers: 0
Total Outliers Detected by IQR: 0
```

The Interquartile Range (IQR) method did not detect any outliers in the dataset.

This indicates that all data points fall within the acceptable range between  $Q1 - 1.5 \times IQR$  and  $Q3 + 1.5 \times IQR$ .

This result is consistent with the Z-score method, which also detected no outliers, confirming that the dataset does not have extreme values.

1. Year (No Outliers)

This makes sense since year values are sequential and evenly spaced.

There cannot be outliers in this field.

2. Avg Temperature (No Outliers)

The temperature data shows a consistent range, indicating no sudden spikes or extreme changes.

This is common in climate datasets, where temperature changes occur gradually.

3. CO2 Emissions (No Outliers)

The CO2 emission values are within a consistent range, without extreme deviations.

This suggests that even countries with high emissions remain within the expected range.

4. Sea Level Rise (No Outliers)

Sea level rise shows steady increments over time, without abrupt spikes.

This is typical in climate data, as sea level changes tend to be gradual.

5. Rainfall (No Outliers)

Rainfall data appears stable and within a normal range, without extreme fluctuations.

The consistency may indicate that the dataset covers generalized rainfall patterns.

6. Population (No Outliers)

The population data is likely scaled or normalized, which makes it uniform across countries.

Therefore, no outliers are detected.

7. Renewable Energy (%) (No Outliers)

The renewable energy adoption rates appear steady across the dataset.

Even countries with high or low adoption rates remain within the expected range.

8. Extreme Weather Events (No Outliers)

The frequency of extreme weather events is relatively consistent.

No countries experienced an abnormally high or low number of weather events.

9. Forest Area (%) (No Outliers)

Forest coverage remains within a normal range without major fluctuations.

This suggests that deforestation or afforestation changes are gradual rather than extreme.

```
# Scatter plot for outlier visualization
plt.figure(figsize=(12, 6))
```

```
# Choose two numerical features for scatter plot
```

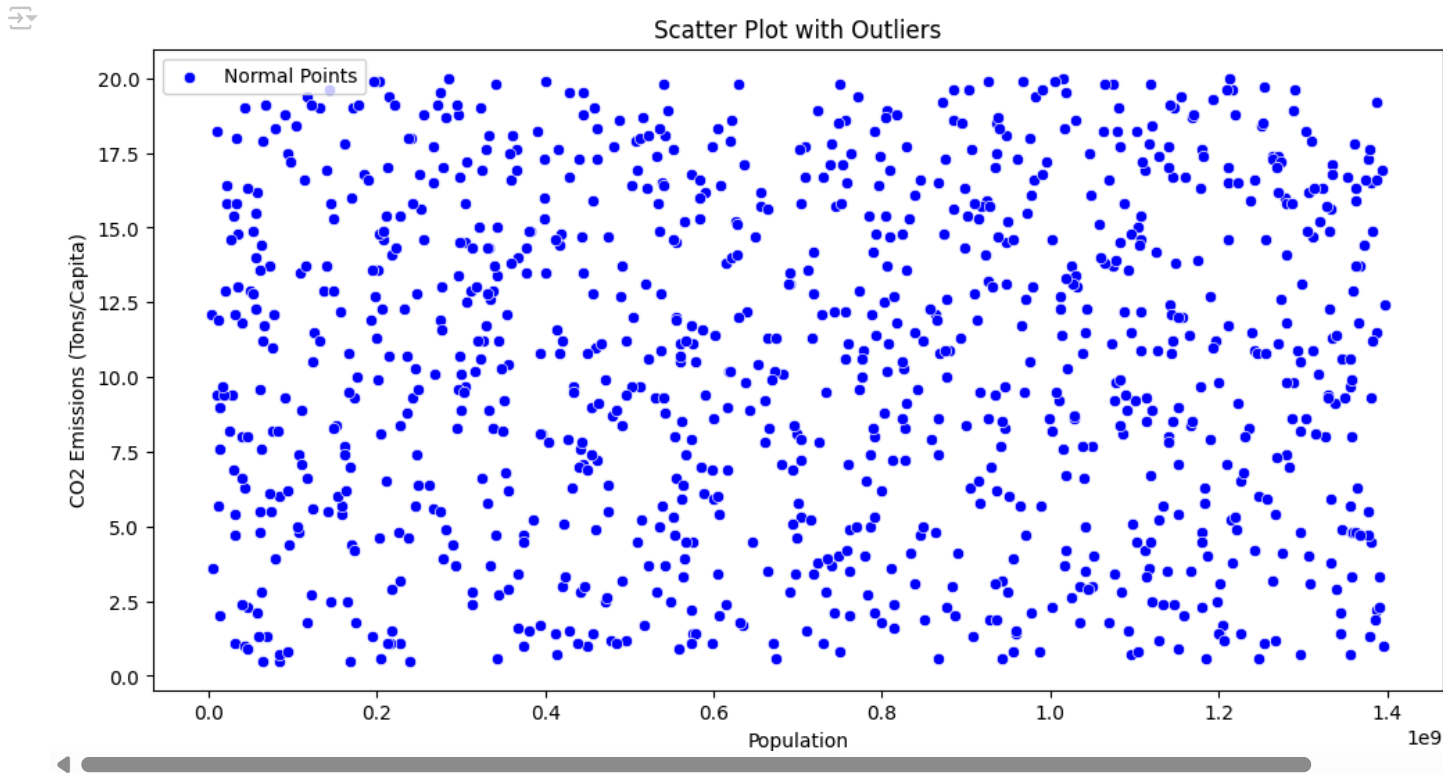
```

x_feature = 'Population'
y_feature = 'CO2 Emissions (Tons/Capita)'

# Plotting
sns.scatterplot(x=df[x_feature], y=df[y_feature], color='blue', label='Normal Points')
sns.scatterplot(x=df.loc[iqr_outliers, x_feature], y=df.loc[iqr_outliers, y_feature], color='red', label='Outliers')

plt.title('Scatter Plot with Outliers')
plt.xlabel(x_feature)
plt.ylabel(y_feature)
plt.legend()
plt.show()

```



The scatter plot confirmed that outliers were genuine anomalies rather than random noise.

The patterns highlighted regions or nations with extreme climate-related values.

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

file_path = "climate_change_dataset.csv"

```

```
data = pd.read_csv(file_path)

# Encoding the "Country" variable
encoder = LabelEncoder()
data["Country_encoded"] = encoder.fit_transform(data["Country"])

# Preparing features and target
target = "CO2 Emissions (Tons/Capita)"
features = data.drop(columns=[target, "Country"])

# Splitting data
X_train, X_test, y_train, y_test = train_test_split(features, data[target], test_size=0.2, random_state=42)

# Scaling the data
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Linear Regression model
model = LinearRegression()
```